



REPORT ON DATA DRIVEN RESEARCH

Alvaro Balderrama

October 3rd, 2024

OUTLINE



- Executive Summary
- 1. Introduction
- 2. Methodology
- 3. Results
 - 3.1. Data collection and data wrangling methodology
 - 3.2. EDA and interactive visual analytics methodology
 - 3.3. Predictive analysis methodology
 - 3.4. EDA with visualization results
 - 3.5. EDA with SQL results
 - 3.6. Interactive map with Folium results
 - 3.7. Plotly Dash dashboard results
 - 3.8. Predictive analysis (classification) results
- 4. Discussion
 - Findings & Implications
- 5. Conclusion
- Appendix

EXECUTIVE SUMMARY



This report was developed in October 2024, as part of the final task of the online course “Data Science and Machine Learning Capstone Project”, part of the Professional Certificate Program: IBM Python Data Science. It is the fifth course of a series which included four other courses: (i) Python Basics for Data Science; (ii) Analyzing Data with Python; (iii) Visualizing Data with Python; (iv) Machine Learning with Python: A Practical Introduction.

The Report includes a summary of the content related to data wrangling, EDA visual analytics, predictive analysis, SQL, interactive map with Folium, Plotly Dash dashboard, and predictive analysis (classification).

1. INTRODUCTION



- Data science is an interdisciplinary field that combines scientific methods, algorithms, and systems to extract knowledge and insights from structured and unstructured data. It involves a blend of skills from areas such as statistics, computer science, and domain expertise, with the goal of uncovering meaningful patterns, trends, and insights that can inform decisions or drive predictions.
- The Professional Certificate program IBM Python Data Science covers the essential aspects needed to kickstart a career that involves data science to some extent.
- The completion of the tasks required presented next, is the final project of a series of courses that started in December 2023 and concluded in October 2024.

2. METHODOLOGY



The process of the courses of this certificate program were documented by taking screenshots during the videos, and compiled in Power Point. For this presentation, some of those slides are presented to summarize specific topics.

A GitHub repository was created for this assignment:

<https://github.com/alvarobalderrama/MLCapstoneProject/tree/main>

3. RESULTS

- 3.1. Data collection and data wrangling methodology
- 3.2. EDA and interactive visual analytics methodology
- 3.3. Predictive analysis methodology
- 3.4. EDA with visualization results
- 3.5. EDA with SQL results
- 3.6. Interactive map with Folium results
- 3.7. Plotly Dash dashboard results
- 3.8. Predictive analysis (classification) results

3.1. Data collection and data wrangling methodology

Before we can continue we must deal with these missing values. The `LandingPad` column will retain None values.

Task 3: Dealing with Missing Values

Calculate below the mean for the `PayloadMass` using the `.mean()`. Then use the mean and the `.replace()`

```
# Calculate the mean value of PayloadMass column
# Replace the np.nan values with its mean value

import numpy as np

# Step 1: Calculate the mean value of the PayloadMass column, ignoring NaN values
mean_payload_mass = data_falcon9['PayloadMass'].mean()

# Step 2: Replace the NaN values in the PayloadMass column with the calculated mean
data_falcon9['PayloadMass'] = data_falcon9['PayloadMass'].replace(np.nan, mean_payload_mass)

# Step 3: Display the updated DataFrame with NaN replaced by the mean
print(data_falcon9.head())
```

✓ 0.0s

	FlightNumber	Date	BoosterVersion	PayloadMass	Orbit	LaunchSite	\
4	1	2010-06-04	Falcon 9	6123.547647	LEO	CCSFS SLC 40	
5	2	2012-05-22	Falcon 9	525.000000	LEO	CCSFS SLC 40	
6	3	2013-03-01	Falcon 9	677.000000	ISS	CCSFS SLC 40	
7	4	2013-09-29	Falcon 9	500.000000	PO	VAFB SLC 4E	
8	5	2013-12-03	Falcon 9	3170.000000	GTO	CCSFS SLC 40	

	Outcome	Flights	GridFins	Reused	Legs	LandingPad	Block	\
4	None None	1	False	False	False	None	1.0	
5	None None	1	False	False	False	None	1.0	
6	None None	1	False	False	False	None	1.0	
7	False Ocean	1	False	False	False	None	1.0	
8	None None	1	False	False	False	None	1.0	

We can use the following line of code to determine the success rate:

```
df["Class"].mean()
✓ 0.0s
0.6666666666666666
```

```
# Define the outcomes that represent a complete failure to land
failure_outcomes = {'False ASDS', 'False Ocean', 'False RTLS', 'None ASDS', 'None None'}

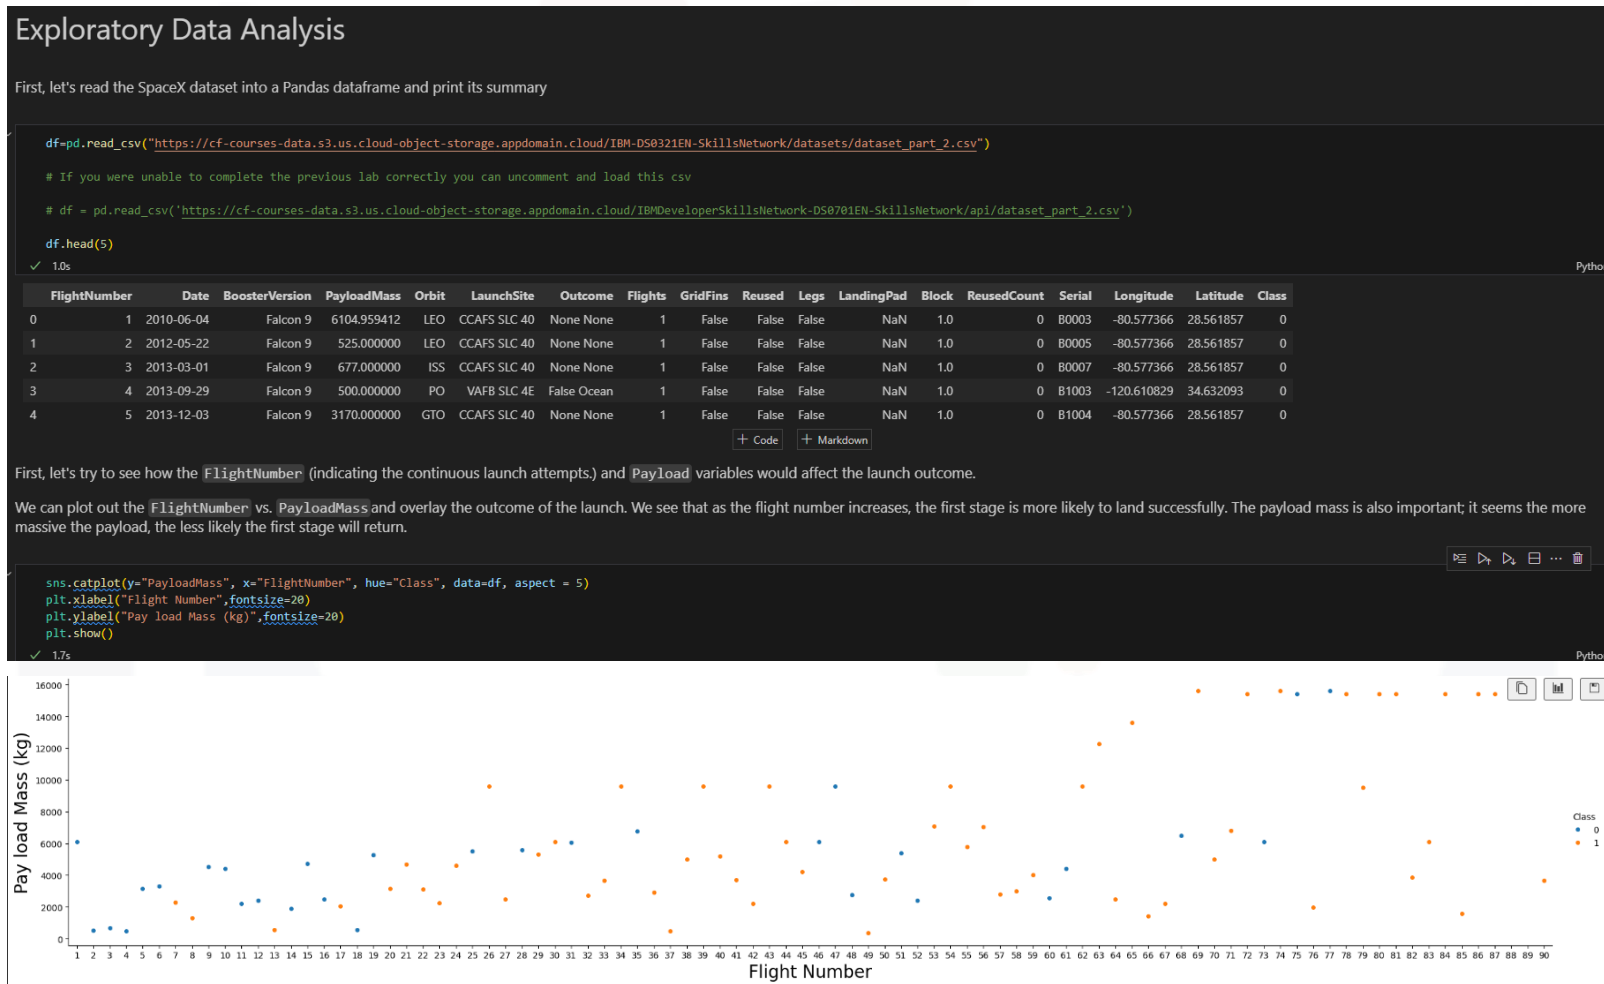
# Count the number of rows in the 'Outcome' column that match the failure outcomes
failure_count = df['Outcome'].isin(failure_outcomes).sum()

# Display the result
print(f"Number of complete failures to land: {failure_count}")
```

✓ 0.0s

Number of complete failures to land: 30

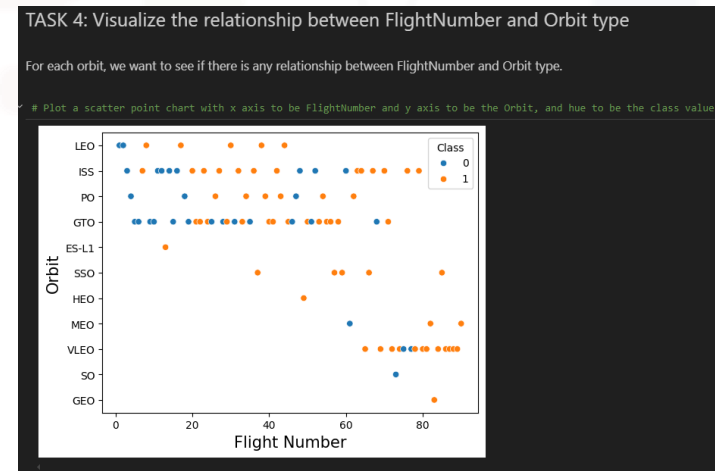
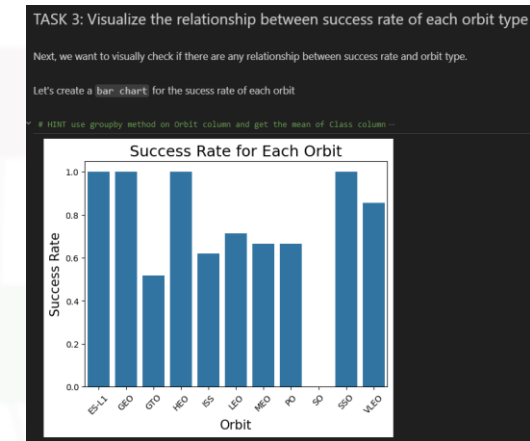
3.2. EDA and interactive visual analytics methodology



3.3. Predictive analysis methodology



3.4. EDA with visualization results



3.5. EDA with SQL results

Task 1

Display the names of the unique launch sites in the space mission

```
%sql SELECT DISTINCT "Launch_Site" FROM SPACEXTBL;
```

```
* sqlite:///my_data1.db  
Done.
```

Launch_Site

CCAFS LC-40

VAFB SLC-4E

KSC LC-39A

CCAFS SLC-40

Task 3

Display the total payload mass carried by boosters launched by NASA (CRS)

+ Code

+ Markdown

```
%sql SELECT SUM("PAYLOAD_MASS_KG_") AS total_payload_mass FROM SPACEXTBL WHERE "Mission_Outcome" LIKE '%CRS%';
```

```
* sqlite:///my_data1.db  
Done.
```

total_payload_mass

None

Task 2

Display 5 records where launch sites begin with the string 'KSC'

```
%sql SELECT * FROM SPACEXTBL WHERE "Launch_Site" LIKE 'KSC%' LIMIT 5;
```

```
* sqlite:///my_data1.db  
Done.
```

Date	Time (UTC)	Booster_Version	Launch_Site	Payload	PAYLOAD_MASS_KG_	Orbit	Customer	Mission_Outcome	Landing_Outcome
2017-02-19	14:39:00	F9 FT B1031.1	KSC LC-39A	SpaceX CRS-10	2490	LEO (ISS)	NASA (CRS)	Success	Success (ground pad)
2017-03-16	6:00:00	F9 FT B1030	KSC LC-39A	EchoStar 23	5600	GTO	EchoStar	Success	No attempt
2017-03-30	22:27:00	F9 FT B1021.2	KSC LC-39A	SES-10	5300	GTO	SES	Success	Success (drone ship)
2017-05-01	11:15:00	F9 FT B1032.1	KSC LC-39A	NROL-76	5300	LEO	NRO	Success	Success (ground pad)
2017-05-15	23:21:00	F9 FT B1034	KSC LC-39A	Inmarsat-5 F4	6070	GTO	Inmarsat	Success	No attempt

Task 10

Rank the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order

```
%sql SELECT "Landing_Outcome", COUNT(*) AS outcome_count FROM SPACEXTBL WHERE "Date" BETWEEN '2010-06-04' AND '2017-03-20' GROUP BY "Landing_Outcome" ORDER BY outcome_count DESC;
```

```
* sqlite:///my_data1.db  
Done.
```

Landing_Outcome	outcome_count
No attempt	10
Success (drone ship)	5
Failure (drone ship)	5
Success (ground pad)	3
Controlled (ocean)	3
Uncontrolled (ocean)	2
Failure (parachute)	2
Precluded (drone ship)	1

3.6. Interactive map with Folium results

Analyze Launch Site Geo Data with Folium

- Mark the locations and proximities of launch sites
- Discover patterns via exploring the map
- Explain how to choose an optimal launch site locations



3.7. Plotly Dash dashboard results

TASK 4: Add a callback function to render the `success-payload-scatter-chart` scatter plot

Next, we want to plot a scatter plot with the x axis to be the payload and the y axis to be the launch outcome (i.e., `class` column). As such, we can visually observe how payload may be correlated with mission outcomes for selected site(s).

In addition, we want to color-label the Booster version on each scatter point so that we may observe mission outcomes with different boosters.

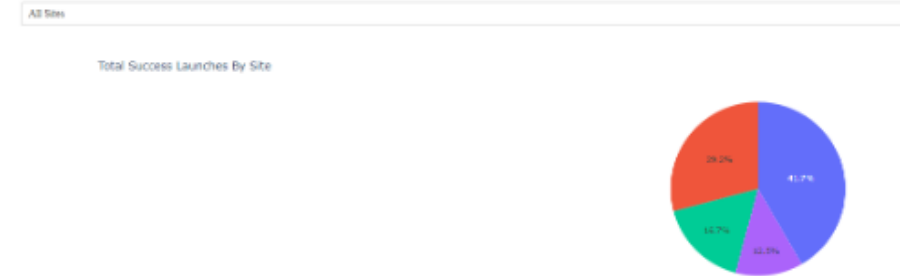
Now, let's add a call function including the following application logic:

- Input to be `[Input(component_id='site-dropdown', component_property='value'), Input(component_id="payload-slider", component_property="value")]`
Note that we have two input components, one to receive selected launch site and another to receive selected payload range
- Output to be `Output(component_id='success-payload-scatter-chart', component_property='figure')`
- A `If-Else` statement to check if ALL sites were selected or just a specific launch site was selected
 - If ALL sites are selected, render a scatter plot to display all values for variable `Payload Mass (kg)` and variable `class`. In addition, the point color needs to be set to the booster version i.e., `color="Booster Version Category"`
 - If a specific launch site is selected, you need to filter the `spacex_df` first, and render a scatter chart to show values `Payload Mass (kg)` and `class` for the selected site, and color-label the point using `Booster Version Category` likewise.

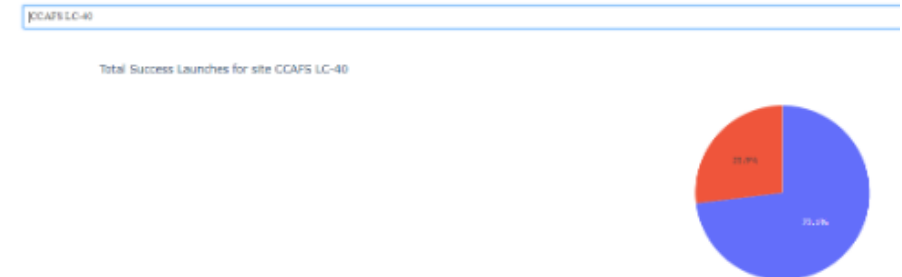
You rendered scatter point should look like the following screenshot:



- Pie chart for all sites are selected



- Pie chart for is selected



```
1 dcc.RangeSlider(id='id',
2                 min=0, max=10000, step=1000,
3                 marks={0: '0',
4                       100: '100'},
5                 value=[min_value, max_value])
```

3.8. Predictive analysis (classification) results

TASK 1

Create a NumPy array from the column `Class` in `data`, by applying the method `to_numpy()` then assign it to the variable `Y`, make sure the output is a Pandas series (only one bracket df['name of column']).

```
# Create a NumPy array from the 'Class' column in the 'data' DataFrame
Y = data['Class'].to_numpy()

# Display the result
print(Y)

[0 0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 1 0 0 1 1 1 1 1 0 1 1 0 1 1 0 1 1 1 1 0 1 1
 1 1 1 1 1 1 1 1 0 0 0 1 1 0 0 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 0 1 1 1 1 1 0 1
 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1]
```

TASK 2

Standardize the data in `<code>X</code>`

 then reassign it to the variable `<code>X</code>` using the transform provided below. ...

```
# students get this ...

from sklearn.preprocessing import StandardScaler ...

[[-1.71291154e+00 -1.94814463e-16 -6.53912840e-01 ... -8.35531692e-01
  1.93309133e+00 -1.93309133e+00]
 [-1.67441914e+00 -1.19523159e+00 -6.53912840e-01 ... -8.35531692e-01
  1.93309133e+00 -1.93309133e+00]
 [-1.63592675e+00 -1.16267307e+00 -6.53912840e-01 ... -8.35531692e-01
  1.93309133e+00 -1.93309133e+00]
 ...
 [ 1.63592675e+00  1.99100483e+00  3.49060516e+00 ...  1.19684269e+00
 -5.17306132e-01  5.17306132e-01]
 [ 1.67441914e+00  1.99100483e+00  1.00389436e+00 ...  1.19684269e+00
 -5.17306132e-01  5.17306132e-01]
 [ 1.71291154e+00 -5.19213966e-01 -6.53912840e-01 ... -8.35531692e-01
 -5.17306132e-01  5.17306132e-01]]
```

TASK 3

Use the function `train_test_split` to split the data `X` and `Y` into training and test data. Set the parameter `test_size` to

```
X_train, X_test, Y_train, Y_test

from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)

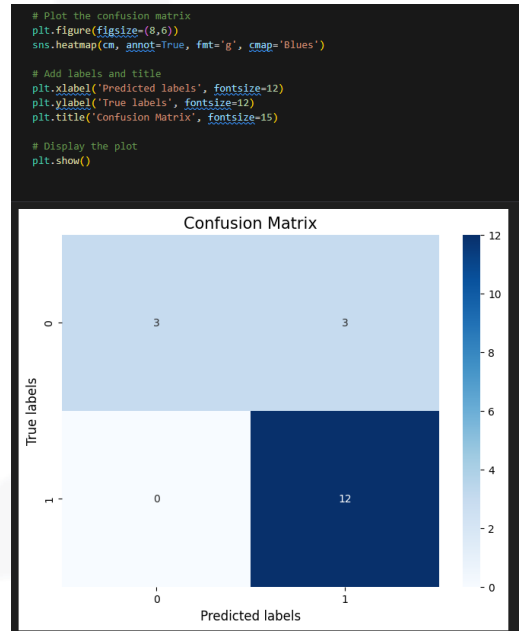
# Display the shapes of the training and testing sets
print("X_train shape: (X_train.shape)")
print("X_test shape: (X_test.shape)")
print("Y_train shape: (Y_train.shape)")
print("Y_test shape: (Y_test.shape)")

X_train shape: (72, 83)
X_test shape: (18, 83)
Y_train shape: (72,)
Y_test shape: (18,)

we can see we only have 18 test samples.

Y_test.shape

(18,)
```



4. DISCUSSION



- The work developed during the Certificate course was presented in the Results section
- The topics covered were data wrangling, EDA, predictive analysis, visualizations, SQL, interactive maps, dashboarding, and classification.

5. CONCLUSION



- Report developed in October 2024 for the final task of the "Data Science and Machine Learning Capstone Project" in the IBM Python Data Science Professional Certificate Program. It is the fifth course after Python Basics, Analyzing Data, Visualizing Data, and Machine Learning with Python.
- Covers data wrangling, EDA, predictive analysis, SQL, Folium maps, Plotly Dash, and classification.
- This report is documented with screenshots from the exercises.
- A GitHub repository was created for the assignment.